

## ■ Machine Learning Interview Questions, Answers & Algorithm Roadmaps

This document provides interview Q&A and **step-by-step roadmaps** for essential machine learning algorithms. Each roadmap explains the algorithm flow, purpose of each step, and basic syntax (no code examples).

---

### ◆ K-Nearest Neighbors (KNN) - Roadmap

**Step 1: Import libraries** Use `from sklearn.neighbors import KNeighborsClassifier`. This loads the algorithm into your workspace.

**Step 2: Feature Scaling** Use `StandardScaler`. KNN uses distance metrics, so features must be on the same scale.

**Step 3: Split data into train/test** Use `train_test_split()` to evaluate generalization.

**Step 4: Fit the model** `model = KNeighborsClassifier(n_neighbors=5)` followed by `model.fit(X_train, y_train)`.

**Step 5: Predict and evaluate** Use `model.predict()` and `accuracy_score()` to check performance.

**Why these steps?** KNN compares data points using distance, so scaling and test splitting are essential. `n_neighbors` (K) should be chosen carefully.

---

### ◆ KMeans Clustering - Roadmap

**Step 1: Preprocessing**

- Remove nulls
- Apply `StandardScaler()`

**Step 2: Choose value of K** Use Elbow method (plot inertia vs. K) to find optimal cluster count.

**Step 3: Initialize and fit** Use `KMeans(n_clusters=K)` and `model.fit(data)`.

**Step 4: Get labels and cluster centers** Use `model.labels_` and `model.cluster_centers_`.

**Step 5: Visualize** Plot using `matplotlib` to understand groupings.

**Why these steps?** Preprocessing ensures equal weight in clustering. Elbow method helps pick effective number of clusters.

---

### ◆ Hierarchical Clustering - Roadmap

**Step 1: Scale features** Use `StandardScaler()` to normalize for fair distance comparison.

**Step 2: Compute linkage matrix** Use `scipy.cluster.hierarchy.linkage()` with `method='ward'` or `'complete'`.

**Step 3: Plot dendrogram** Use `dendrogram()` to visualize data hierarchy.

**Step 4: Choose number of clusters** Cut dendrogram at a specific height (threshold) to form clusters.

**Step 5: Assign labels** Use `fcluster()` to label each data point.

**Why these steps?** Hierarchical clustering shows how data clusters step-by-step, and dendrograms help in visual decision-making.

---

### ◆ Apriori Algorithm - Roadmap

**Step 1: Transaction Data Preprocessing**

- Convert dataset to boolean using `df.applymap(lambda x: 1 if x > 0 else 0)`

**Step 2: Generate frequent itemsets** Use `apriori()` from `mlxtend.frequent_patterns` with `min_support`.

**Step 3: Generate association rules** Use `association_rules()` with `metric='confidence'` and `min_threshold`.

**Step 4: Filter strong rules** Use `lift > 1` and `confidence > 0.6` for meaningful patterns.

**Why these steps?** Apriori finds itemsets frequently occurring together. Data must be in binary form to evaluate presence/absence of items.

---

### ◆ Linear Regression - Roadmap

**Step 1: Handle missing values & outliers** Linear regression is sensitive to outliers. Clean your dataset.

**Step 2: Check linear relationship** Use scatter plots or correlation matrix.

**Step 3: Split data** Use `train_test_split()`.

**Step 4: Fit the model** Use `LinearRegression()` and `.fit(X_train, y_train)`.

**Step 5: Evaluate** Check  $R^2$  score, MSE, or RMSE using metrics module.

**Why these steps?** Linear regression assumes a linear pattern and is affected by data quality. Evaluation confirms model quality.

---

#### ◆ Logistic Regression - Roadmap

**Step 1: Check class imbalance** Use `value_counts()` to verify if one class dominates.

**Step 2: Apply scaling** Improves convergence.

**Step 3: Fit the model** Use `LogisticRegression()` and `.fit()`.

**Step 4: Predict and evaluate** Use classification report, confusion matrix, ROC-AUC.

**Step 5: Apply regularization** Use `penalty='l1'` or `'l2'` to reduce overfitting.

**Why these steps?** Logistic regression outputs probabilities and works best with balanced, scaled data. Regularization improves generalization.

---

#### ◆ Decision Trees - Roadmap

**Step 1: Preprocess data** Trees don't require scaling but handle missing or categorical values properly.

**Step 2: Fit model** Use `DecisionTreeClassifier()` with `criterion='gini'` or `'entropy'`.

**Step 3: Tune hyperparameters** Control depth, `min_samples_split`, etc., to avoid overfitting.

**Step 4: Evaluate model** Use accuracy, confusion matrix, and tree visualization.

**Why these steps?** Tree models learn rules, but without constraints can overfit. Tuning is necessary.

---

#### ◆ Random Forest - Roadmap

**Step 1: Preprocess** No scaling required. Handle nulls and categorical variables.

**Step 2: Fit model** Use `RandomForestClassifier()` and fit on training data.

**Step 3: Evaluate using OOB score** Set `oob_score=True` to get out-of-bag error estimate.

**Step 4: Feature importance** Use `model.feature_importances_` to interpret.

**Step 5: Hyperparameter tuning** Use `GridSearchCV` or `RandomizedSearchCV`.

**Why these steps?** Random Forest reduces overfitting via ensemble and gives useful insight into feature relevance.

---

### ◆ Support Vector Machine (SVM) - Roadmap

**Step 1: Scale the data** SVM is very sensitive to feature scale. Use StandardScaler.

**Step 2: Fit model** Use SVC() with kernel='linear' or 'rbf'. Fit on training data.

**Step 3: Tune parameters** C (regularization) and gamma (kernel coefficient) using GridSearch.

**Step 4: Evaluate performance** Use confusion matrix, accuracy, and ROC-AUC.

**Step 5: Visualize margin (for 2D)** Use support vectors and decision boundary plots.

**Why these steps?** SVM works best when data is scaled and hyperparameters are tuned.

---

### ◆ Naive Bayes - Roadmap

**Step 1: Format categorical/text data** Use label encoding for categorical or CountVectorizer/TfidfVectorizer for text.

**Step 2: Choose model** Use GaussianNB, MultinomialNB, or BernoulliNB based on data type.

**Step 3: Fit the model** Use .fit(X\_train, y\_train).

**Step 4: Predict and evaluate** Use .predict() and classification metrics.

**Step 5: Analyze prior and likelihood** Use .class\_prior\_ and .theta\_ for understanding model internals.

**Why these steps?** Naive Bayes is fast and suitable for text but assumes feature independence, which may limit performance on some data.